

CS 2710 Foundations of AI

Lecture 9

Propositional logic

Milos Hauskrecht

milos@cs.pitt.edu


5329 Sennott Square

CS 2710 Foundations of AI

Knowledge representation

CS 2710 Foundations of AI

Knowledge-based agent



- **Knowledge base (KB):**
 - A set of sentences that describe facts about the world in some formal (representational) language
 - **Domain specific**
- **Inference engine:**
 - A set of procedures that use the representational language to infer new facts from known ones or answer a variety of KB queries. Inferences typically require search.
 - **Domain independent**

Example: MYCIN

- MYCIN: an expert system for diagnosis of bacterial infections
- **Knowledge base** represents
 - Facts about a specific patient case
 - Rules describing relations between entities in the bacterial infection domain

If	1. The stain of the organism is gram-positive, and 2. The morphology of the organism is coccus, and 3. The growth conformation of the organism is chains
Then	the identity of the organism is streptococcus

- **Inference engine:**
 - manipulates the facts and known relations to answer diagnostic queries (consistent with findings and rules)

Knowledge representation

- The objective of knowledge representation is to express the knowledge about the world in a computer-tractable form
- Key aspects of knowledge representation languages:
 - **Syntax**: describes how sentences are formed in the language
 - **Semantics**: describes the meaning of sentences, what is it the sentence refers to in the real world
 - **Computational aspect**: describes how sentences and objects are manipulated in concordance with semantical conventions

Many KB systems rely on some variant of logic

CS 2710 Foundations of AI

Logic

A formal language for expressing knowledge and ways of reasoning.

Logic is defined by:

- **A set of sentences**
 - A sentence is constructed from a set of primitives according to syntax rules.
- **A set of interpretations**
 - An interpretation gives a semantic to primitives. It associates primitives with values.
- **The valuation (meaning) function V**
 - Assigns a value (typically the truth value) to a given sentence under some interpretation
$$V : \text{sentence} \times \text{interpretation} \rightarrow \{True, False\}$$

CS 2710 Foundations of AI

Example of logic

Language of numerical constraints:

- **A sentence:**

$$x + 3 \leq z$$

x, z - variable symbols (primitives in the language)

- **An interpretation:**

$$I: x = 5, z = 2$$

Variables mapped to specific real numbers

- **Valuation (meaning) function V :**

$V(x + 3 \leq z, I)$ is **False** for $I: x = 5, z = 2$

is **True** for $I: x = 5, z = 10$

Types of logic

- Different types of logics possible:

- Propositional logic
- First-order logic
- Temporal logic
- Numerical constraints logic
- Map-coloring logic

In the following:

- **Propositional logic.**

- Formal language for making logical inferences
- Foundations of **propositional logic**: **George Boole** (1854)

Propositional logic. Syntax

- **Propositional logic P:**
 - defines a language for symbolic reasoning
- **Proposition:** a statement that is either true or false
- Examples of propositions:
 - *Pitt is located in the Oakland section of Pittsburgh.*
 - *France is in Europe.*
 - *It rains outside.*
 - *2 is a prime number and 6 is a prime*
 - *How are you?* Not a proposition.

CS 2710 Foundations of AI

Propositional logic. Syntax

- **Formally propositional logic P:**
 - Is defined by **Syntax+interpretation+semantics of P**

Syntax:

- **Symbols (alphabet)** in P:
 - **Constants:** *True, False*
 - **Propositional symbols**
Examples:
 - *P*
 - *Pitt is located in the Oakland section of Pittsburgh.,*
 - *It rains outside,* etc.
 - **A set of connectives:**

$\neg, \wedge, \vee, \Rightarrow, \Leftrightarrow$

CS 2710 Foundations of AI

Propositional logic. Syntax

Sentences in the propositional logic:

- **Atomic sentences:**
 - **Constructed from constants and propositional symbols**
 - True, False are (atomic) sentences
 - P, Q or *Light in the room is on*, *It rains outside* are (atomic) sentences
- **Composite sentences:**
 - **Constructed from valid sentences via connectives**
 - If A, B are sentences then
$$\neg A \quad (A \wedge B) \quad (A \vee B) \quad (A \Rightarrow B) \quad (A \Leftrightarrow B)$$
or $(A \vee B) \wedge (A \vee \neg B)$ are sentences

Propositional logic. Semantics.

The semantic gives the meaning to sentences.

the semantics in the propositional logic is defined by:

- 1. Interpretation of propositional symbols and constants**
 - Semantics of atomic sentences
- 2. Through the meaning of connectives**
 - Meaning (semantics) of composite sentences

Semantic: propositional symbols

A **propositional symbol**

- a statement about the world that is either true or false

Examples:

- *Pitt is located in the Oakland section of Pittsburgh*
- *It rains outside*
- *Light in the room is on*

- An **interpretation** maps symbols to one of the two values: **True (T)**, or **False (F)**, depending on whether the symbol is satisfied in the world

I: *Light in the room is on* -> **True**, *It rains outside* -> **False**

I': *Light in the room is on* -> **False**, *It rains outside* -> **False**

Semantic: propositional symbols

The **meaning (value)** of the propositional symbol for a specific interpretation is given by its interpretation

I: *Light in the room is on* -> **True**, *It rains outside* -> **False**

$V(\text{Light in the room is on}, \mathbf{I}) = \text{True}$

$V(\text{It rains outside}, \mathbf{I}) = \text{False}$

I': *Light in the room is on* -> **False**, *It rains outside* -> **False**

$V(\text{Light in the room is on}, \mathbf{I}') = \text{False}$

Semantics: constants

- The meaning (truth) of constants:

- True and False constants are always (under any interpretation) assigned the corresponding *True, False* value

$$\left. \begin{array}{l} V(\text{True}, \mathbf{I}) = \text{True} \\ V(\text{False}, \mathbf{I}) = \text{False} \end{array} \right\} \text{For any interpretation } \mathbf{I}$$

Semantics: composite sentences.

- The meaning (truth value) of complex propositional sentences.

- Determined using the standard rules of logic:

P	Q	$\neg P$	$P \wedge Q$	$P \vee Q$	$P \Rightarrow Q$	$P \Leftrightarrow Q$
<i>True</i>	<i>True</i>	<i>False</i>	<i>True</i>	<i>True</i>	<i>True</i>	<i>True</i>
<i>True</i>	<i>False</i>	<i>False</i>	<i>False</i>	<i>True</i>	<i>False</i>	<i>False</i>
<i>False</i>	<i>True</i>	<i>True</i>	<i>False</i>	<i>True</i>	<i>True</i>	<i>False</i>
<i>False</i>	<i>False</i>	<i>True</i>	<i>False</i>	<i>False</i>	<i>True</i>	<i>True</i>

Translation

Assume the following sentences:

- It is not sunny this afternoon and it is colder than yesterday.
- We will go swimming only if it is sunny.
- If we do not go swimming then we will take a canoe trip.
- If we take a canoe trip, then we will be home by sunset.

Denote:

- p = It is sunny this afternoon
- q = it is colder than yesterday
- r = We will go swimming
- s = we will take a canoe trip
- t = We will be home by sunset

Translation

Assume the following sentences:

- It is not sunny this afternoon and it is colder than yesterday. $\neg p \wedge q$
- We will go swimming only if it is sunny.
- If we do not go swimming then we will take a canoe trip.
- If we take a canoe trip, then we will be home by sunset.

Denote:

- p = It is sunny this afternoon
- q = it is colder than yesterday
- r = We will go swimming
- s = we will take a canoe trip
- t = We will be home by sunset

Translation

Assume the following sentences:

- It is not sunny this afternoon and it is colder than yesterday. $\neg p \wedge q$
- We will go swimming only if it is sunny. $r \rightarrow p$
- If we do not go swimming then we will take a canoe trip.
- If we take a canoe trip, then we will be home by sunset.

Denote:

- p = It is sunny this afternoon
- q = it is colder than yesterday
- r = We will go swimming
- s = we will take a canoe trip
- t = We will be home by sunset

Translation

Assume the following sentences:

- It is not sunny this afternoon and it is colder than yesterday. $\neg p \wedge q$
- We will go swimming only if it is sunny. $r \rightarrow p$
- If we do not go swimming then we will take a canoe trip. $\neg r \rightarrow s$
- If we take a canoe trip, then we will be home by sunset.

Denote:

- p = It is sunny this afternoon
- q = it is colder than yesterday
- r = We will go swimming
- s = we will take a canoe trip
- t = We will be home by sunset

Translation

Assume the following sentences:

- It is not sunny this afternoon and it is colder than yesterday. $\neg p \wedge q$
- We will go swimming only if it is sunny. $r \rightarrow p$
- If we do not go swimming then we will take a canoe trip. $\neg r \rightarrow s$
- If we take a canoe trip, then we will be home by sunset. $s \rightarrow t$

Denote:

- p = It is sunny this afternoon
- q = it is colder than yesterday
- r = We will go swimming
- s = we will take a canoe trip
- t = We will be home by sunset

Contradiction and Tautology

Some composite sentences may always (under any interpretation) evaluate to a single truth value:

- **Contradiction** (always *False*)

$$P \wedge \neg P$$

- **Tautology** (always *True*)

$$P \vee \neg P$$

$$\left. \begin{array}{l} \neg(P \vee Q) \Leftrightarrow (\neg P \wedge \neg Q) \\ \neg(P \wedge Q) \Leftrightarrow (\neg P \vee \neg Q) \end{array} \right\} \text{DeMorgan's Laws}$$

Model, validity and satisfiability

- A **model (in logic)**: An interpretation is a model for a set of sentences if it assigns true to each sentence in the set.
- A sentence is **satisfiable** if it has a model;
 - There is at least one interpretation under which the sentence can evaluate to True.
- A sentence is **valid** if it is **True** in all interpretations
 - i.e., if its negation is **not satisfiable** (leads to contradiction)

P	Q	$P \vee Q$	$(P \vee Q) \wedge \neg Q$	$((P \vee Q) \wedge \neg Q) \Rightarrow P$
<i>True</i>	<i>True</i>	<i>True</i>	<i>False</i>	<i>True</i>
<i>True</i>	<i>False</i>	<i>True</i>	<i>True</i>	<i>True</i>
<i>False</i>	<i>True</i>	<i>True</i>	<i>False</i>	<i>True</i>
<i>False</i>	<i>False</i>	<i>False</i>	<i>False</i>	<i>True</i>

CS 2710 Foundations of AI

Model, validity and satisfiability

- A **model (in logic)**: An interpretation is a model for a set of sentences if it assigns true to each sentence in the set.
- A sentence is **satisfiable** if it has a model;
 - There is at least one interpretation under which the sentence can evaluate to True.
- A sentence is **valid** if it is **True** in all interpretations
 - i.e., if its negation is **not satisfiable** (leads to contradiction)

P	Q	$P \vee Q$	$(P \vee Q) \wedge \neg Q$	$((P \vee Q) \wedge \neg Q) \Rightarrow P$
<i>True</i>	<i>True</i>	<i>True</i>	<i>False</i>	<i>True</i>
<i>True</i>	<i>False</i>	<i>True</i>	<i>True</i>	<i>True</i>
<i>False</i>	<i>True</i>	<i>True</i>	<i>False</i>	<i>True</i>
<i>False</i>	<i>False</i>	<i>False</i>	<i>False</i>	<i>True</i>

CS 2710 Foundations of AI

Model, validity and satisfiability

- A **model (in logic)**: An interpretation is a model for a set of sentences if it assigns true to each sentence in the set.
- A sentence is **satisfiable** if it has a model;
 - There is at least one interpretation under which the sentence can evaluate to True.
- A sentence is **valid** if it is **True** in all interpretations
 - i.e., if its negation is **not satisfiable** (leads to contradiction)

Satisfiable sentence

P	Q	$P \vee Q$	$(P \vee Q) \wedge \neg Q$	$((P \vee Q) \wedge \neg Q) \Rightarrow P$
<i>True</i>	<i>True</i>	<i>True</i>	<i>False</i>	<i>True</i>
<i>True</i>	<i>False</i>	<i>True</i>	<i>True</i>	<i>True</i>
<i>False</i>	<i>True</i>	<i>True</i>	<i>False</i>	<i>True</i>
<i>False</i>	<i>False</i>	<i>False</i>	<i>False</i>	<i>True</i>

Model, validity and satisfiability

- A **model (in logic)**: An interpretation is a model for a set of sentences if it assigns true to each sentence in the set.
- A sentence is **satisfiable** if it has a model;
 - There is at least one interpretation under which the sentence can evaluate to True.
- A sentence is **valid** if it is **True** in all interpretations
 - i.e., if its negation is **not satisfiable** (leads to contradiction)


Satisfiable sentence

Valid sentence

P	Q	$P \vee Q$	$(P \vee Q) \wedge \neg Q$	$((P \vee Q) \wedge \neg Q) \Rightarrow P$
<i>True</i>	<i>True</i>	<i>True</i>	<i>False</i>	<i>True</i>
<i>True</i>	<i>False</i>	<i>True</i>	<i>True</i>	<i>True</i>
<i>False</i>	<i>True</i>	<i>True</i>	<i>False</i>	<i>True</i>
<i>False</i>	<i>False</i>	<i>False</i>	<i>False</i>	<i>True</i>

Entailment

- **Entailment** reflects the relation of one fact in the world following from the others



- Entailment $KB \models \alpha$
- Knowledge base KB entails sentence α if and only if α is true in all worlds where KB is true

Sound and complete inference.

Inference is a process by which conclusions are reached.

- We want to implement the inference process on a computer !!

Assume an **inference procedure** i that

- derives a sentence α from the KB : $KB \vdash_i \alpha$

Properties of the inference procedure in terms of entailment

- **Soundness:** An inference procedure is **sound**

If $KB \vdash_i \alpha$ then it is true that $KB \models \alpha$

- **Completeness:** An inference procedure is **complete**

If $KB \models \alpha$ then it is true that $KB \vdash_i \alpha$

Logical inference problem

Logical inference problem:

- **Given:**
 - a knowledge base KB (a set of sentences) and
 - a sentence α (called **a theorem**),
- **Does a KB semantically entail α ?** $KB \models \alpha$?

In other words: In all interpretations in which sentences in the KB are true, is also α true?

Question: Is there a procedure (program) that can decide this problem in a finite number of steps?

Answer: Yes. Logical inference problem for the propositional logic is **decidable**.

Solving logical inference problem

In the following:

How to design the procedure that answers:

$$KB \models \alpha ?$$

Three approaches:

- **Truth-table approach**
- **Inference rules**
- **Conversion to the inverse SAT problem**
 - **Resolution-refutation**

Truth-table approach

Problem: $KB \models \alpha$?

- We need to check all possible interpretations for which the KB is true (models of KB) whether α is true for each of them

Truth table:

- enumerates truth values of sentences for all possible interpretations (assignments of True/False values to propositional symbols)

Example:

		KB		α
P	Q	$P \vee Q$	$P \Leftrightarrow Q$	$(P \vee \neg Q) \wedge Q$
True	True	True	True	True
True	False	True	False	False
False	True	True	False	False
False	False	False	True	False

CS 2710 Foundations of AI

Truth-table approach

Problem: $KB \models \alpha$?

- We need to check all possible interpretations for which the KB is true (models of KB) whether α is true for each of them

Truth table:

- enumerates truth values of sentences for all possible interpretations (assignments of True/False to propositional symbols)

Example:

		KB		α
P	Q	$P \vee Q$	$P \Leftrightarrow Q$	$(P \vee \neg Q) \wedge Q$
True	True	True	True	True
True	False	True	False	False
False	True	True	False	False
False	False	False	True	False

CS 2710 Foundations of AI

Truth-table approach

Problem: $KB \models \alpha$?

- We need to check all possible interpretations for which the KB is true (models of KB) whether α is true for each of them

Truth table:

- enumerates truth values of sentences for all possible interpretations (assignments of True/False to propositional symbols)

Example:

		KB		α
P	Q	$P \vee Q$	$P \Leftrightarrow Q$	$(P \vee \neg Q) \wedge Q$
True	True	True	True	True
True	False	True	False	False
False	True	True	False	False
False	False	False	True	False



Truth-table approach

A two steps procedure:

1. **Generate table for all possible interpretations**
2. Check whether the sentence α evaluates to true whenever KB evaluates to true

Example: $KB = (A \vee C) \wedge (B \vee \neg C)$ $\alpha = (A \vee B)$

A	B	C	$A \vee C$	$(B \vee \neg C)$	KB	α
True	True	True				
True	True	False				
True	False	True				
True	False	False				
False	True	True				
False	True	False				
False	False	True				
False	False	False				

Truth-table approach

A two steps procedure:

1. Generate table for all possible interpretations
2. Check whether the sentence α evaluates to true whenever KB evaluates to true

Example: $KB = (A \vee C) \wedge (B \vee \neg C)$ $\alpha = (A \vee B)$

<i>A</i>	<i>B</i>	<i>C</i>	$A \vee C$	$(B \vee \neg C)$	KB	α
<i>True</i>	<i>True</i>	<i>True</i>	<i>True</i>	<i>True</i>	<i>True</i>	<i>True</i>
<i>True</i>	<i>True</i>	<i>False</i>	<i>True</i>	<i>True</i>	<i>True</i>	<i>True</i>
<i>True</i>	<i>False</i>	<i>True</i>	<i>True</i>	<i>False</i>	<i>False</i>	<i>True</i>
<i>True</i>	<i>False</i>	<i>False</i>	<i>True</i>	<i>True</i>	<i>True</i>	<i>True</i>
<i>False</i>	<i>True</i>	<i>True</i>	<i>True</i>	<i>True</i>	<i>True</i>	<i>True</i>
<i>False</i>	<i>True</i>	<i>False</i>	<i>False</i>	<i>True</i>	<i>False</i>	<i>True</i>
<i>False</i>	<i>False</i>	<i>True</i>	<i>True</i>	<i>False</i>	<i>False</i>	<i>False</i>
<i>False</i>	<i>False</i>	<i>False</i>	<i>False</i>	<i>True</i>	<i>False</i>	<i>False</i>

CS 2710 Foundations of AI

Truth-table approach

A two steps procedure:

1. Generate table for all possible interpretations
2. Check whether the sentence α evaluates to true whenever KB evaluates to true

Example: $KB = (A \vee C) \wedge (B \vee \neg C)$ $\alpha = (A \vee B)$

<i>A</i>	<i>B</i>	<i>C</i>	$A \vee C$	$(B \vee \neg C)$	KB	α
<i>True</i>	<i>True</i>	<i>True</i>	<i>True</i>	<i>True</i>	<i>True</i>	<i>True</i>
<i>True</i>	<i>True</i>	<i>False</i>	<i>True</i>	<i>True</i>	<i>True</i>	<i>True</i>
<i>True</i>	<i>False</i>	<i>True</i>	<i>True</i>	<i>False</i>	<i>False</i>	<i>True</i>
<i>True</i>	<i>False</i>	<i>False</i>	<i>True</i>	<i>True</i>	<i>True</i>	<i>True</i>
<i>False</i>	<i>True</i>	<i>True</i>	<i>True</i>	<i>True</i>	<i>True</i>	<i>True</i>
<i>False</i>	<i>True</i>	<i>False</i>	<i>False</i>	<i>True</i>	<i>False</i>	<i>True</i>
<i>False</i>	<i>False</i>	<i>True</i>	<i>True</i>	<i>False</i>	<i>False</i>	<i>False</i>
<i>False</i>	<i>False</i>	<i>False</i>	<i>False</i>	<i>True</i>	<i>False</i>	<i>False</i>



CS 2710 Foundations of AI

Truth-table approach

$$KB = (A \vee C) \wedge (B \vee \neg C) \quad \alpha = (A \vee B)$$

<i>A</i>	<i>B</i>	<i>C</i>	<i>A</i> \vee <i>C</i>	(<i>B</i> \vee \neg <i>C</i>)	<i>KB</i>	α
<i>True</i>	<i>True</i>	<i>True</i>	<i>True</i>	<i>True</i>	<i>True</i>	<i>True</i>
<i>True</i>	<i>True</i>	<i>False</i>	<i>True</i>	<i>True</i>	<i>True</i>	<i>True</i>
<i>True</i>	<i>False</i>	<i>True</i>	<i>True</i>	<i>False</i>	<i>False</i>	<i>True</i>
<i>True</i>	<i>False</i>	<i>False</i>	<i>True</i>	<i>True</i>	<i>True</i>	<i>True</i>
<i>False</i>	<i>True</i>	<i>True</i>	<i>True</i>	<i>True</i>	<i>True</i>	<i>True</i>
<i>False</i>	<i>True</i>	<i>False</i>	<i>False</i>	<i>True</i>	<i>False</i>	<i>True</i>
<i>False</i>	<i>False</i>	<i>True</i>	<i>True</i>	<i>False</i>	<i>False</i>	<i>False</i>
<i>False</i>	<i>False</i>	<i>False</i>	<i>False</i>	<i>True</i>	<i>False</i>	<i>False</i>

KB entails α

- The **truth-table approach** is **sound and complete** for the propositional logic!!

CS 2710 Foundations of AI

Limitations of the truth table approach.

$$KB \models \alpha ?$$

What is the computational complexity of the truth table approach?

- ?

CS 2710 Foundations of AI

Limitations of the truth table approach.

$$KB \models \alpha ?$$

What is the computational complexity of the truth table approach?

Exponential in the number of the proposition symbols

2^n Rows in the table has to be filled

Limitations of the truth table approach.

$$KB \models \alpha ?$$

What is the computational complexity of the truth table approach?

Exponential in the number of the proposition symbols

2^n Rows in the table has to be filled

But typically only for a small subset of rows the KB is true

Limitations of the truth table approach.

$$KB \models \alpha ?$$

Problem with the truth table approach:

- the truth table is **exponential** in the number of propositional symbols (we checked all assignments)
- KB is true on only a smaller subset

Limitation of the truth table approach.

$$KB \models \alpha ?$$

Problem with the truth table approach:

- the truth table is **exponential** in the number of propositional symbols (we checked all assignments)
- KB is true only on a small subset interpretations

How to make the process more efficient?

Inference rules approach.

$$KB \models \alpha ?$$

Problem with the truth table approach:

- the truth table is **exponential** in the number of propositional symbols (we checked all assignments)
- KB is true on only a smaller subset

How to make the process more efficient?

Solution: check only entries for which KB is *True*.

This is the idea behind the inference rules approach



Inference rules:

- Represent sound inference patterns repeated in inferences
- Can be used to generate new (sound) sentences from the existing ones

Inference rules for logic

• Modus ponens

$$\frac{A \Rightarrow B, \quad A}{B}$$

 premise
 conclusion

- If both sentences in the premise are true then conclusion is true.
- The modus ponens inference rule is **sound**.
 - We can prove this through the truth table.

<i>A</i>	<i>B</i>	<i>A</i> \Rightarrow <i>B</i>
<i>False</i>	<i>False</i>	<i>True</i>
<i>False</i>	<i>True</i>	<i>True</i>
<i>True</i>	<i>False</i>	<i>False</i>
<i>True</i>	<i>True</i>	<i>True</i>

Inference rules for logic

- And-elimination

$$\frac{A_1 \wedge A_2 \wedge \dots \wedge A_n}{A_i}$$

- And-introduction

$$\frac{A_1, A_2, \dots, A_n}{A_1 \wedge A_2 \wedge \dots \wedge A_n}$$

- Or-introduction

$$\frac{A_i}{A_1 \vee A_2 \vee \dots \vee A_i \vee \dots \vee A_n}$$

Inference rules for logic

- Elimination of double negation

$$\frac{\neg\neg A}{A}$$


- Unit resolution

$$\frac{A \vee B, \quad \neg A}{B}$$

- Resolution

$$\frac{A \vee B, \quad \neg B \vee C}{A \vee C}$$

A special
case of



- All of the above inference rules **are sound**. We can prove this through the truth table, similarly to the **modus ponens** case.

Example. Inference rules approach.

KB: $P \wedge Q$ $P \Rightarrow R$ $(Q \wedge R) \Rightarrow S$ **Theorem:** S

1. $P \wedge Q$
2. $P \Rightarrow R$
3. $(Q \wedge R) \Rightarrow S$

Example. Inference rules approach.

KB: $P \wedge Q$ $P \Rightarrow R$ $(Q \wedge R) \Rightarrow S$ **Theorem:** S

1. $P \wedge Q$
2. $P \Rightarrow R$
3. $(Q \wedge R) \Rightarrow S$
4. P

From 1 and And-elim

$$\frac{A_1 \wedge A_2 \wedge \dots \wedge A_n}{A_i}$$

Example. Inference rules approach.

KB: $P \wedge Q$ $P \Rightarrow R$ $(Q \wedge R) \Rightarrow S$ **Theorem:** S

1. $P \wedge Q$
2. $P \Rightarrow R$
3. $(Q \wedge R) \Rightarrow S$
4. P
5. R

From 2,4 and Modus ponens

$$\frac{A \Rightarrow B, \quad A}{B}$$

Example. Inference rules approach.

KB: $P \wedge Q$ $P \Rightarrow R$ $(Q \wedge R) \Rightarrow S$ **Theorem:** S

1. $P \wedge Q$
2. $P \Rightarrow R$
3. $(Q \wedge R) \Rightarrow S$
4. P
5. R
6. Q

From 1 and And-elim

$$\frac{A_1 \wedge A_2 \wedge \quad A_n}{A_i}$$

Example. Inference rules approach.

KB: $P \wedge Q$ $P \Rightarrow R$ $(Q \wedge R) \Rightarrow S$ **Theorem:** S

1. $P \wedge Q$
2. $P \Rightarrow R$
3. $(Q \wedge R) \Rightarrow S$
4. P
5. R
6. Q
7. $(Q \wedge R)$

From 5,6 and And-introduction

$$\frac{A_1, A_2, \quad A_n}{A_1 \wedge A_2 \wedge \quad A_n}$$

Example. Inference rules approach.

KB: $P \wedge Q$ $P \Rightarrow R$ $(Q \wedge R) \Rightarrow S$ **Theorem:** S

1. $P \wedge Q$
2. $P \Rightarrow R$
3. $(Q \wedge R) \Rightarrow S$
4. P
5. R
6. Q
7. $(Q \wedge R)$
8. S

$$\frac{A \Rightarrow B, \quad A}{B}$$

From 7,3 and Modus ponens

Proved: S

Example. Inference rules approach.

KB: $P \wedge Q$ $P \Rightarrow R$ $(Q \wedge R) \Rightarrow S$ **Theorem:** S

1. $P \wedge Q$
2. $P \Rightarrow R$
3. $(Q \wedge R) \Rightarrow S$
4. P From 1 and And-elim
5. R From 2,4 and Modus ponens
6. Q From 1 and And-elim
7. $(Q \wedge R)$ From 5,6 and And-introduction
8. S From 7,3 and Modus ponens

Proved: S


CS 2710 Foundations of AI

Inference rules

- To show that theorem α holds for a KB
 - we may need to apply a number of sound inference rules

Problem: many possible inference rules to be applied next

Looks familiar?




CS 2710 Foundations of AI

Logic inferences and search

- To show that theorem α holds for a KB
 - we may need to apply a number of sound inference rules

Problem: many possible rules to can be applied next

Looks familiar?



This is an instance of a search problem:

Truth table method (from the search perspective):

- blind enumeration and checking

Logic inferences and search

Inference rule method as a search problem:

- **State:** a set of sentences that are known to be true
- **Initial state:** a set of sentences in the KB
- **Operators:** applications of inference rules
 - Allow us to add new sound sentences to old ones
- **Goal state:** a theorem α is derived from KB

Logic inference:

- **Proof:** A sequence of sentences that are immediate consequences of applied inference rules
- **Theorem proving:** process of finding a proof of theorem